

```
In [1]: import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
import seaborn as sns
from lightgbm import LGBMClassifier
import gc
import matplotlib.font_manager as fm
from sklearn.metrics import roc_auc_score
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import train_test_split
plt.rcParams['font.sans-serif'] = ['SimHei'] # 中文字体设置
plt.rcParams['axes.unicode_minus'] = False # 解决保存图像是负号 '-' 显示为方块的问题
sns.set(font='SimHei') # 解决Seaborn中文显示问题
from sklearn import preprocessing
from sklearn.preprocessing import LabelEncoder
```

由于CSV文件内数据超过4千万笔，Pandas DataFrame无法一次读取这么大的数据量，因此将CSV文件内的数据每100万笔读取一次，然后从这100万笔中随机抽取5%的数据，形成一个新的训练集

```
In [2]: chunksize = 10 ** 6
num_of_chunk = 0
train = pd.DataFrame()

for chunk in pd.read_csv('train.csv', chunksize=chunksize):
    num_of_chunk += 1
    train = pd.concat([train, chunk.sample(frac=.05, replace=False, random_state=37)], axis=0)
    print('Processing Chunk No. ' + str(num_of_chunk))

train.reset_index(inplace=True)

# 备份原始训练集的数据长度，方便后续重新分割索引
train_len = len(train)
train_len
```

```
Processing Chunk No. 1
Processing Chunk No. 2
Processing Chunk No. 3
Processing Chunk No. 4
Processing Chunk No. 5
Processing Chunk No. 6
Processing Chunk No. 7
Processing Chunk No. 8
Processing Chunk No. 9
Processing Chunk No. 10
Processing Chunk No. 11
Processing Chunk No. 12
Processing Chunk No. 13
Processing Chunk No. 14
Processing Chunk No. 15
Processing Chunk No. 16
Processing Chunk No. 17
Processing Chunk No. 18
Processing Chunk No. 19
Processing Chunk No. 20
Processing Chunk No. 21
Processing Chunk No. 22
Processing Chunk No. 23
Processing Chunk No. 24
Processing Chunk No. 25
Processing Chunk No. 26
Processing Chunk No. 27
Processing Chunk No. 28
Processing Chunk No. 29
Processing Chunk No. 30
Processing Chunk No. 31
Processing Chunk No. 32
Processing Chunk No. 33
Processing Chunk No. 34
Processing Chunk No. 35
Processing Chunk No. 36
Processing Chunk No. 37
Processing Chunk No. 38
Processing Chunk No. 39
Processing Chunk No. 40
Processing Chunk No. 41
```

Out[2]: 2021448

将test测试集的资料读出后，把train训练集和 test测试集合并成新的数据集df，方便同时对资料集进行数据预处理。

```
In [3]: ▶ test_file = 'test.gz'
df = pd.concat([train, pd.read_csv(test_file, compression='gzip')]).drop(['index', 'id'], axis=1)
```

本次预测的目的为用10天的点击情况作为训练模型，用来预测第11天的点击情况，因此年月日期没有参考价值，但是weekday有参考价值。所以把原始数据中hour特征中的日期，转化为weekday。每天的小时段根据每人不同的生活习惯，有参考价值。为避免24小时的纬度导致数据崩溃，采用十二地支计时法将小时的区间分为12个时辰，每两小时一个时辰。

```
In [4]: ▶ # 定义一个将hour数据转换为日期格式的函数
def get_date(hour):
    y = '20'+str(hour)[:2]
    m = str(hour)[2:4]
    d = str(hour)[4:6]
    return y+'-'+m+'-'+d

# 创建weekday列，将hour数据转换后填入
df['weekday'] = pd.to_datetime(df.hour.apply(get_date)).dt.dayofweek.astype(str)

# 定义一个将hour数据转换为12时辰的函数
def tran_hour(x):
    x = x % 100
    while x in [23,0]:
        return '1'
    while x in [1,2]:
        return '2'
    while x in [3,4]:
        return '3'
    while x in [5,6]:
        return '4'
    while x in [7,8]:
        return '5'
    while x in [9,10]:
        return '6'
    while x in [11,12]:
        return '7'
    while x in [13,14]:
        return '8'
    while x in [15,16]:
        return '9'
    while x in [17,18]:
        return '10'
    while x in [19,20]:
        return '1'
    while x in [21,22]:
        return '12'

# 将hour列的数据转换为时段
df['hour'] = df.hour.apply(tran_hour)
```

In [5]:

查看数据集的的简要摘要
df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 6598912 entries, 0 to 4577463
Data columns (total 24 columns):
Column Dtype
--- -
0 click float64
1 hour object
2 C1 int64
3 banner_pos int64
4 site_id object
5 site_domain object
6 site_category object
7 app_id object
8 app_domain object
9 app_category object
10 device_id object
11 device_ip object
12 device_model object
13 device_type int64
14 device_conn_type int64
15 C14 int64
16 C15 int64
17 C16 int64
18 C17 int64
19 C18 int64
20 C19 int64
21 C20 int64
22 C21 int64
23 weekday object
dtypes: float64(1), int64(12), object(11)
memory usage: 1.2+ GB

In [6]:

df.head(20)

Out[6]:

	click	hour	C1	banner_pos	site_id	site_domain	site_category	app_id	app_domain	app_category	...	device_conn_type
0	1.0	3	1005	0	1fbe01fe	f3845767	28905ebd	ecad2386	7801e8d9	07d7df22	...	0 1
1	0.0	3	1005	0	1fbe01fe	f3845767	28905ebd	ecad2386	7801e8d9	07d7df22	...	0 1
2	0.0	3	1005	0	1fbe01fe	f3845767	28905ebd	ecad2386	7801e8d9	07d7df22	...	0 1
3	0.0	4	1005	0	1fbe01fe	f3845767	28905ebd	ecad2386	7801e8d9	07d7df22	...	0 1
4	0.0	1	1005	0	6256f5b4	28f93029	f028772b	ecad2386	7801e8d9	07d7df22	...	0 2
5	0.0	2	1005	1	d9750ee7	98572c79	f028772b	ecad2386	7801e8d9	07d7df22	...	0 1
6	0.0	3	1005	0	78d60190	1b32ed33	70fb0e29	ecad2386	7801e8d9	07d7df22	...	0 2
7	1.0	2	1005	0	1fbe01fe	f3845767	28905ebd	ecad2386	7801e8d9	07d7df22	...	0 1
8	0.0	2	1005	0	85f751fd	c4e18dd6	50e219e0	39947756	2347f47a	cef3e649	...	2 2
9	0.0	3	1005	0	85f751fd	c4e18dd6	50e219e0	febd1138	82e27996	0f2161f8	...	0 2
10	0.0	3	1005	0	ebfa440f	4973dc87	f028772b	ecad2386	7801e8d9	07d7df22	...	0 2
11	0.0	4	1005	0	85f751fd	c4e18dd6	50e219e0	e96773f0	2347f47a	0f2161f8	...	0 2
12	0.0	3	1005	0	686546d3	9d54950b	f028772b	ecad2386	7801e8d9	07d7df22	...	0 2
13	0.0	4	1005	0	85f751fd	c4e18dd6	50e219e0	7e7baafa	2347f47a	0f2161f8	...	0 2
14	1.0	2	1005	0	1fbe01fe	f3845767	28905ebd	ecad2386	7801e8d9	07d7df22	...	0 1
15	0.0	3	1010	1	85f751fd	c4e18dd6	50e219e0	8c0dcd5a	7801e8d9	0f2161f8	...	0 2
16	0.0	3	1005	0	1fbe01fe	f3845767	28905ebd	ecad2386	7801e8d9	07d7df22	...	0 1
17	0.0	3	1005	0	85f751fd	c4e18dd6	50e219e0	54c5d545	2347f47a	0f2161f8	...	0 2
18	0.0	3	1005	0	85f751fd	c4e18dd6	50e219e0	19aaa345	d9b5648e	0f2161f8	...	0 2
19	1.0	1	1005	0	1fbe01fe	f3845767	28905ebd	ecad2386	7801e8d9	07d7df22	...	0 1

20 rows × 24 columns

```
In [7]: df.isnull().any() #查看缺失值情况
```

```
Out[7]: click                True
hour                  False
C1                   False
banner_pos           False
site_id              False
site_domain          False
site_category        False
app_id               False
app_domain           False
app_category         False
device_id            False
device_ip            False
device_model         False
device_type          False
device_conn_type     False
C14                  False
C15                  False
C16                  False
C17                  False
C18                  False
C19                  False
C20                  False
C21                  False
weekday              False
dtype: bool
```

确认每个特征的计数值，int类别的特征，最多包含4333个特征值，在这个新生成的600万笔数据的资料集，既然不是连续计数，据此判定本资料集的所有特征均为Object类型的数值。

```
In [8]: len_of_feature_count = []
for i in df.columns[2:23].tolist():
    print(i, ': ', len(df[i].astype(str).value_counts()))
    len_of_feature_count.append(len(df[i].astype(str).value_counts()))
```

```
C1 : 7
banner_pos : 7
site_id : 3501
site_domain : 4304
site_category : 25
app_id : 5456
app_domain : 309
app_category : 29
device_id : 553513
device_ip : 1876695
device_model : 6302
device_type : 4
device_conn_type : 4
C14 : 2662
C15 : 8
C16 : 9
C17 : 467
C18 : 4
C19 : 68
C20 : 167
C21 : 62
```

```
In [9]: df.columns
```

```
Out[9]: Index(['click', 'hour', 'C1', 'banner_pos', 'site_id', 'site_domain',
              'site_category', 'app_id', 'app_domain', 'app_category', 'device_id',
              'device_ip', 'device_model', 'device_type', 'device_conn_type', 'C14',
              'C15', 'C16', 'C17', 'C18', 'C19', 'C20', 'C21', 'weekday'],
              dtype='object')
```

```
In [10]: df.info()


<class 'pandas.core.frame.DataFrame'>
Int64Index: 6598912 entries, 0 to 4577463
Data columns (total 24 columns):
#   Column          Dtype
---  -
0   click           float64
1   hour            object
2   C1              int64
3   banner_pos      int64
4   site_id         object
5   site_domain     object
6   site_category   object
7   app_id          object
8   app_domain      object
9   app_category    object
10  device_id       object
11  device_ip       object
12  device_model    object
13  device_type     int64
14  device_conn_type int64
15  C14             int64
16  C15             int64
17  C16             int64
18  C17             int64
19  C18             int64
20  C19             int64
21  C20             int64
22  C21             int64
23  weekday         object
dtypes: float64(1), int64(12), object(11)
memory usage: 1.2+ GB
```

```
In [11]: df_describe = df.describe()
df_describe
```

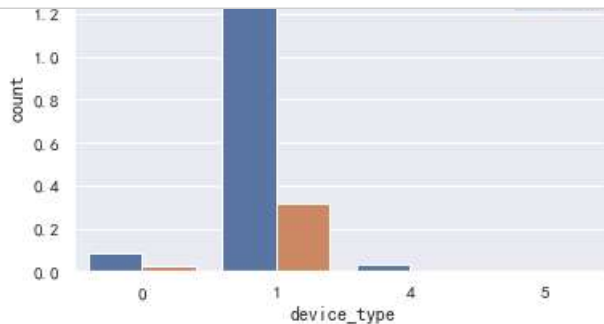
Out[11]:

	click	C1	banner_pos	device_type	device_conn_type	C14	C15	C16	C17
count	2.021448e+06	6.598912e+06	6.598912e+06	6.598912e+06	6.598912e+06	6.598912e+06	6.598912e+06	6.598912e+06	6.598912e+06
mean	1.699584e-01	1.004982e+03	2.345505e-01	1.011268e+00	3.615288e-01	2.083368e+04	3.198177e+02	5.972008e+01	2.359957e+01
std	3.755963e-01	9.838387e-01	4.839890e-01	4.606424e-01	8.839961e-01	4.374423e+03	2.664230e+01	4.676491e+01	5.457326e+01
min	0.000000e+00	1.001000e+03	0.000000e+00	0.000000e+00	0.000000e+00	3.750000e+02	1.200000e+02	2.000000e+01	1.120000e+01
25%	0.000000e+00	1.005000e+03	0.000000e+00	1.000000e+00	0.000000e+00	1.977200e+04	3.200000e+02	5.000000e+01	2.227000e+01
50%	0.000000e+00	1.005000e+03	0.000000e+00	1.000000e+00	0.000000e+00	2.226100e+04	3.200000e+02	5.000000e+01	2.545000e+01
75%	0.000000e+00	1.005000e+03	0.000000e+00	1.000000e+00	0.000000e+00	2.372500e+04	3.200000e+02	5.000000e+01	2.707000e+01
max	1.000000e+00	1.012000e+03	7.000000e+00	5.000000e+00	5.000000e+00	2.434900e+04	1.024000e+03	1.024000e+03	2.793000e+01

```
In [12]: need_tran_feature = ['site_id', 'site_domain', 'site_category', 'app_id', 'app_domain', 'app_category', 'device_ip', 'device_id']
# 依次将非int类型的数据转换为int类型
df['weekday'] = df['weekday'].astype(str).astype(int)
df['hour'] = df['hour'].astype(str).astype(int)
for column in need_tran_feature:
    enc = LabelEncoder()
    df[column] = enc.fit_transform(df[column])
#热力图
sub_df = df.iloc[:, 3:]
plt.figure(figsize=(20, 18))
sns.heatmap(sub_df.corr().abs(), annot=True)
plt.show()
```



```
In [13]: # 确认所有特征的分布情况
plt_features=['click', 'hour', 'C1', 'banner_pos', 'site_id', 'site_domain',
             'site_category', 'app_id', 'app_domain', 'app_category', 'device_model', 'device_type', 'device_conn_type', 'C14',
             'C15', 'C16', 'C17', 'C18', 'C19', 'C20', 'C21', 'weekday']
for i in plt_features:
    sns.countplot(x = i, hue = "click", data = df)
    plt.show()
```



```
In [14]: # 根据处理过的df资料集，重新将train训练集，test测试集分割出来
train = df[:train_len]
test = df[train_len:]
```

由于资料集的数据量非常的庞大，同时正向label的占比仅占全部资料的17%左右，比例明显失衡，需要用强化加权功能的演算法，因此决定使用微软的lightgbm演算法，解决强化权重问题

```
In [15]: x=gc.collect()#释放内存
# 从train训练集中，标签为0的资料中，随机抽取与标签为1一样多的数量，并将其结合成正反标签各占50%的资料集
pre_X = train[train['click'] == 0].sample(n=len(train[train['click'] == 1]), random_state=37)
pre_X = pd.concat([pre_X, train[train['click'] == 1]]).sample(frac=1)
pre_y = pre_X[['click']]
pre_X.drop(['click'], axis=1, inplace=True)
test.drop(['click'], axis=1, inplace=True)
```

C:\Anaconda\lib\site-packages\pandas\core\frame.py:3990: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
return super().drop()

```
In [16]: # 将新的资料集分割为训练集和验证集
X_train, X_test, y_train, y_test = train_test_split(pre_X, pre_y, test_size=0.20, stratify=pre_y, random_state=37)
```

```
In [17]: clf = LGBMClassifier(  
    n_estimators=2000,  
    learning_rate=0.02,  
    subsample=0.8,  
    colsample_bytree=0.4,  
    metric='auc',  
    num_leaves=20,  
    )  
h = clf.fit(X_train, y_train, eval_set=[(X_test, y_test)], verbose=50, early_stopping_rounds=100)
```

C:\Anaconda\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
    return f(*args, **kwargs)
```

Training until validation scores don't improve for 100 rounds


```
[50]   valid_0's auc: 0.702218  
[100]  valid_0's auc: 0.709815  
[150]  valid_0's auc: 0.715355  
[200]  valid_0's auc: 0.719456  
[250]  valid_0's auc: 0.722052  
[300]  valid_0's auc: 0.724263  
[350]  valid_0's auc: 0.726016  
[400]  valid_0's auc: 0.727612  
[450]  valid_0's auc: 0.728888  
[500]  valid_0's auc: 0.730108  
[550]  valid_0's auc: 0.7312  
[600]  valid_0's auc: 0.732072  
[650]  valid_0's auc: 0.732867  
[700]  valid_0's auc: 0.733585  
[750]  valid_0's auc: 0.734235  
[800]  valid_0's auc: 0.734829  
[850]  valid_0's auc: 0.735453  
[900]  valid_0's auc: 0.73597  
[950]  valid_0's auc: 0.736412  
[1000] valid_0's auc: 0.736799  
[1050] valid_0's auc: 0.737153  
[1100] valid_0's auc: 0.737488  
[1150] valid_0's auc: 0.737789  
[1200] valid_0's auc: 0.738101  
[1250] valid_0's auc: 0.738396  
[1300] valid_0's auc: 0.738669  
[1350] valid_0's auc: 0.738937  
[1400] valid_0's auc: 0.739172  
[1450] valid_0's auc: 0.739455  
[1500] valid_0's auc: 0.739666  
[1550] valid_0's auc: 0.73991  
[1600] valid_0's auc: 0.74014  
[1650] valid_0's auc: 0.740382  
[1700] valid_0's auc: 0.740563  
[1750] valid_0's auc: 0.740731  
[1800] valid_0's auc: 0.740915  
[1850] valid_0's auc: 0.741068  
[1900] valid_0's auc: 0.741291  
[1950] valid_0's auc: 0.741417  
[2000] valid_0's auc: 0.741568
```

Did not meet early stopping. Best iteration is:

```
[2000] valid_0's auc: 0.741568
```



```
In [18]: feature_imp = pd.DataFrame(sorted(zip(clf.feature_importances_, X_train.columns)), columns=['Value', 'Feature'])
plt.figure(figsize=(20, 10))
#sns.barplot(x="Value", y="Feature", data=feature_imp.sort_values(by="Value", ascending=False).iloc[:50])
sns.barplot(x="Value", y="Feature", data=feature_imp.sort_values(by="Value", ascending=False))
plt.title("重要特征排序")
plt.tight_layout()
plt.show()
```



```
In [19]: features = X_train.columns

feature_importance_values = clf.feature_importances_

feature_importances = pd.DataFrame({'feature': list(features), 'importance': feature_importance_values})

feature_importances.sort_values('importance', inplace=True, ascending=False)
print(feature_importances)
```

	feature	importance
4	site_domain	4161
3	site_id	4025
6	app_id	3653
11	device_model	3113
14	C14	2939
17	C17	2119
19	C19	1860
21	C21	1859
20	C20	1842
9	device_id	1426
10	device_ip	1424
7	app_domain	1388
5	site_category	1295
8	app_category	1107
0	hour	1020
22	weekday	811
18	C18	725
2	banner_pos	701
1	C1	688
13	device_conn_type	606
16	C16	561
12	device_type	347
15	C15	330

```
In [20]: y_pred = clf.predict(X_train)
#y_pred = clf.predict(X)
print("Roc_auc_score: ", roc_auc_score(y_train, y_pred)*100, "%")


# 输出混淆矩阵，查看预测结果
confmat = confusion_matrix(y_true=y_train, y_pred=y_pred, labels=[0, 1])

fig, ax = plt.subplots(figsize=(2.5, 2.5))
ax.matshow(confmat, cmap=plt.cm.Blues, alpha=0.3)
for i in range(confmat.shape[0]):
    for j in range(confmat.shape[1]):
        ax.text(x=j, y=i, s=confmat[i, j], va='center', ha='center')

plt.xlabel('Predicted label')
plt.ylabel('True label')

plt.tight_layout()
plt.show()
```

Roc_auc_score: 68.2839245583395 %



```
In [21]: x=gc.collect()#释放内存
```